

## content of the current directory:

In [1]: % ls

```
Readme.md                      winequality-red.csv
Wine Quality Analysis.ipynb    winequality-white.csv
Wine+Quality+Analysis.html     winequality.names.txt
```

In [2]:

```
# import libs
%matplotlib inline
import pandas as pd
import seaborn as sns; sns.set(style="whitegrid", palette="muted")
import numpy as np
import matplotlib.pyplot as plt
```

## Create DataFrames for white and red wines

In [3]:

```
white_wine_df = pd.read_csv('winequality-white.csv', sep=";")
red_wine_df = pd.read_csv('winequality-red.csv', sep=";")
```

## DataFrames for red and white wines combined

In [4]:

```
ww = white_wine_df.loc[:]
ww["color"] = "white"
rw = red_wine_df.loc[:]
rw["color"] = "red"
wine_df = pd.concat([ww, rw], ignore_index=True)
```

## Data

In [5]: white\_wine\_df.head()

Out[5]:

|   | fixed acidity | volatile acidity | citric acid | residual sugar | chlorides | free sulfur dioxide | total sulfur dioxide | density | pH   | sulphates | a |
|---|---------------|------------------|-------------|----------------|-----------|---------------------|----------------------|---------|------|-----------|---|
| 0 | 7.0           | 0.27             | 0.36        | 20.7           | 0.045     | 45.0                | 170.0                | 1.0010  | 3.00 | 0.45      | 8 |
| 1 | 6.3           | 0.30             | 0.34        | 1.6            | 0.049     | 14.0                | 132.0                | 0.9940  | 3.30 | 0.49      | 9 |
| 2 | 8.1           | 0.28             | 0.40        | 6.9            | 0.050     | 30.0                | 97.0                 | 0.9951  | 3.26 | 0.44      | 1 |
| 3 | 7.2           | 0.23             | 0.32        | 8.5            | 0.058     | 47.0                | 186.0                | 0.9956  | 3.19 | 0.40      | 9 |
| 4 | 7.2           | 0.23             | 0.32        | 8.5            | 0.058     | 47.0                | 186.0                | 0.9956  | 3.19 | 0.40      | 9 |

```
In [6]: red_wine_df.head()
```

Out[6]:

|   | fixed acidity | volatile acidity | citric acid | residual sugar | chlorides | free sulfur dioxide | total sulfur dioxide | density | pH   | sulphates | quality |
|---|---------------|------------------|-------------|----------------|-----------|---------------------|----------------------|---------|------|-----------|---------|
| 0 | 7.4           | 0.70             | 0.00        | 1.9            | 0.076     | 11.0                | 34.0                 | 0.9978  | 3.51 | 0.56      | 9       |
| 1 | 7.8           | 0.88             | 0.00        | 2.6            | 0.098     | 25.0                | 67.0                 | 0.9968  | 3.20 | 0.68      | 9       |
| 2 | 7.8           | 0.76             | 0.04        | 2.3            | 0.092     | 15.0                | 54.0                 | 0.9970  | 3.26 | 0.65      | 9       |
| 3 | 11.2          | 0.28             | 0.56        | 1.9            | 0.075     | 17.0                | 60.0                 | 0.9980  | 3.16 | 0.58      | 9       |
| 4 | 7.4           | 0.70             | 0.00        | 1.9            | 0.076     | 11.0                | 34.0                 | 0.9978  | 3.51 | 0.56      | 9       |

```
In [7]: assert white_wine_df.columns.all() == red_wine_df.columns.all()  
        ", ".join(list(white_wine_df.columns))
```

```
Out[7]: 'fixed acidity,volatile acidity,citric acid,residual sugar,chlorides,free sulfur dioxide,total sulfur dioxide,density,pH,sulphates,alcohol,quality,color'
```

## test for null values and check correct datatypes

```
In [8]: assert white_wine_df.notnull().all().all()  
        white_wine_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 4898 entries, 0 to 4897  
Data columns (total 13 columns):  
fixed acidity      4898 non-null float64  
volatile acidity   4898 non-null float64  
citric acid        4898 non-null float64  
residual sugar     4898 non-null float64  
chlorides          4898 non-null float64  
free sulfur dioxide 4898 non-null float64  
total sulfur dioxide 4898 non-null float64  
density           4898 non-null float64  
pH                4898 non-null float64  
sulphates         4898 non-null float64  
alcohol           4898 non-null float64  
quality           4898 non-null int64  
color             4898 non-null object  
dtypes: float64(11), int64(1), object(1)  
memory usage: 497.5+ KB
```

no null values in white wine dataframe found

```
In [9]: assert red_wine_df.notnull().all().all()  
red_wine_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 1599 entries, 0 to 1598  
Data columns (total 13 columns):  
fixed acidity      1599 non-null float64  
volatile acidity   1599 non-null float64  
citric acid        1599 non-null float64  
residual sugar     1599 non-null float64  
chlorides          1599 non-null float64  
free sulfur dioxide 1599 non-null float64  
total sulfur dioxide 1599 non-null float64  
density           1599 non-null float64  
pH                1599 non-null float64  
sulphates         1599 non-null float64  
alcohol           1599 non-null float64  
quality           1599 non-null int64  
color             1599 non-null object  
dtypes: float64(11), int64(1), object(1)  
memory usage: 162.5+ KB
```

no null values in red wine dataframe found

All datatypes are numeric.

## Build categoricals

```
In [10]: # this can be crucial :)  
white_wine_df["color"] = white_wine_df["color"].astype("category")  
red_wine_df["color"] = red_wine_df["color"].astype("category")  
wine_df["color"] = wine_df["color"].astype("category")
```

## Means

### White Wines:

```
In [11]: white_wine_df.mean()
```

```
Out[11]: fixed acidity      6.854788  
volatile acidity   0.278241  
citric acid        0.334192  
residual sugar     6.391415  
chlorides          0.045772  
free sulfur dioxide 35.308085  
total sulfur dioxide 138.360657  
density           0.994027  
pH                3.188267  
sulphates         0.489847  
alcohol           10.514267  
quality           5.877909  
dtype: float64
```

## Red Wines:

```
In [12]: red_wine_df.mean()
```

```
Out[12]: fixed acidity      8.319637  
volatile acidity    0.527821  
citric acid         0.270976  
residual sugar      2.538806  
chlorides           0.087467  
free sulfur dioxide 15.874922  
total sulfur dioxide 46.467792  
density             0.996747  
pH                  3.311113  
sulphates           0.658149  
alcohol             10.422983  
quality             5.636023  
dtype: float64
```

## Differences between red and white wine means that are greater than 1.0

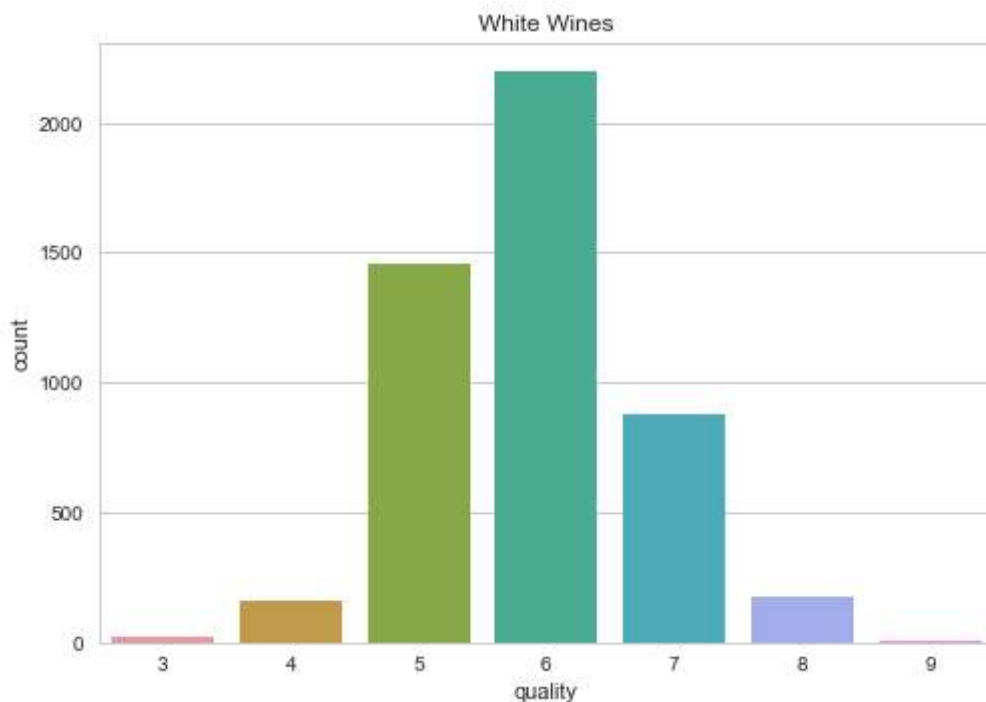
```
In [13]: mean_diff = white_wine_df.mean() - red_wine_df.mean()  
mean_diff_abs = mean_diff.apply(lambda x: abs(x))  
mean_diff[mean_diff_abs >= 1.0]
```

```
Out[13]: fixed acidity      -1.464850  
residual sugar             3.852609  
free sulfur dioxide        19.433163  
total sulfur dioxide       91.892865  
dtype: float64
```

## Distribution of Quality

```
In [14]: sns.countplot(data=white_wine_df, x="quality")  
sns.plt.title("White Wines")
```

```
Out[14]: <matplotlib.text.Text at 0x1112b1ac8>
```

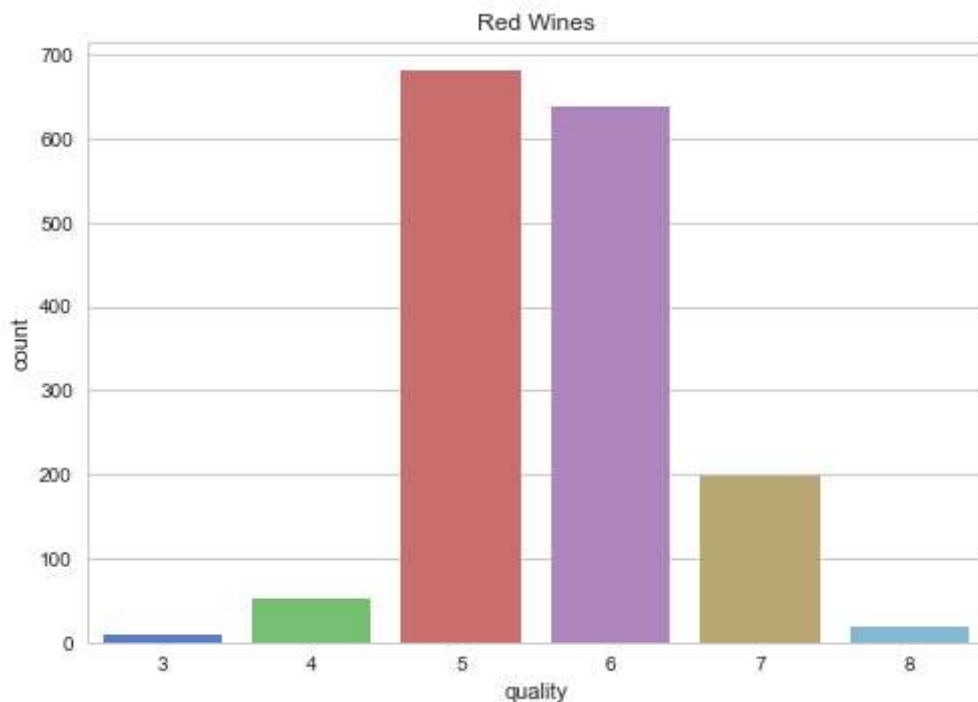


```
In [15]: white_wine_df.quality.describe()
```

```
Out[15]: count      4898.000000  
mean         5.877909  
std          0.885639  
min          3.000000  
25%          5.000000  
50%          6.000000  
75%          6.000000  
max          9.000000  
Name: quality, dtype: float64
```

```
In [16]: x = sns.countplot(data=red_wine_df, x="quality")  
sns.plt.title("Red Wines")
```

```
Out[16]: <matplotlib.text.Text at 0x11144de80>
```



```
In [17]: red_wine_df.quality.describe()
```

```
Out[17]: count      1599.000000  
mean         5.636023  
std          0.807569  
min          3.000000  
25%          5.000000  
50%          6.000000  
75%          6.000000  
max          8.000000  
Name: quality, dtype: float64
```

## What may be important for a high quality rating?

To find out, the percentual mean differences for low quality to high quality wines over the total mean are calculated, resulting in percentual changes.

**For white wines:**

```
In [18]: x = white_wine_df.groupby(["quality"]).mean()
lower_qual = x.loc[:4].mean()
higher_qual = x.loc[7:].mean()
ww_perc_means = (higher_qual - lower_qual) / white_wine_df.mean() * 100
ww_perc_means
```

```
Out[18]: alcohol      14.068993
chlorides      -38.372538
citric acid       7.758984
density      -0.254610
fixed acidity    -6.235608
free sulfur dioxide -10.177344
pH              1.934262
quality              NaN
residual sugar   -8.100321
sulphates        1.999155
total sulfur dioxide -18.439304
volatile acidity -27.979143
dtype: float64
```

**Comparing low quality means to high quality ones, the following attributes differ more than 5 per cent:**

```
In [19]: ww_perc_means[abs(ww_perc_means) > 5]
```

```
Out[19]: alcohol      14.068993
chlorides      -38.372538
citric acid       7.758984
fixed acidity    -6.235608
free sulfur dioxide -10.177344
residual sugar   -8.100321
total sulfur dioxide -18.439304
volatile acidity -27.979143
dtype: float64
```

**Comparing low quality means to high quality ones, the following attributes differ more than 10 per cent:**

```
In [20]: ww_perc_means[abs(ww_perc_means) > 10]
```

```
Out[20]: alcohol      14.068993
chlorides      -38.372538
free sulfur dioxide -10.177344
total sulfur dioxide -18.439304
volatile acidity -27.979143
dtype: float64
```

**For red wines:**

```
In [21]: x = red_wine_df.groupby(["quality"]).mean()
lower_qual = x.loc[:4].mean()
higher_qual = x.loc[7:].mean()
rw_perc_means = (higher_qual - lower_qual) / red_wine_df.mean() * 100
rw_perc_means
```

```
Out[21]: alcohol          16.023546
chlorides             -38.955960
citric acid           77.707371
density              -0.134937
fixed acidity          7.811538
free sulfur dioxide   12.783852
pH                   -3.345302
quality                NaN
residual sugar        -0.609713
sulphates             26.028988
total sulfur dioxide   7.875629
volatile acidity      -71.161438
dtype: float64
```

**Comparing low quality means to high quality ones, the following attributes differ more than 5 per cent:**

```
In [22]: rw_perc_means[abs(rw_perc_means) > 5]
```

```
Out[22]: alcohol          16.023546
chlorides             -38.955960
citric acid           77.707371
fixed acidity          7.811538
free sulfur dioxide   12.783852
sulphates             26.028988
total sulfur dioxide   7.875629
volatile acidity      -71.161438
dtype: float64
```

**Comparing low quality means to high quality ones, the following attributes differ more than 10 per cent:**

```
In [23]: rw_perc_means[abs(rw_perc_means) > 10]
```

```
Out[23]: alcohol          16.023546
chlorides             -38.955960
citric acid           77.707371
free sulfur dioxide   12.783852
sulphates             26.028988
volatile acidity      -71.161438
dtype: float64
```

**What will be taken a closer look at:**


- Alcohol
- Chlorides
- Citric Acid
- Sulphates
- Sulfur Dioxides
- Volatile Acidity

# Sulfur Dioxides and Quality

## White Wines

```
In [24]: fig, (ax1, ax2) = plt.subplots(1,2)
fig.set_size_inches(14.5, 4.5)
fig.dpi = 300
sns.stripplot(data=white_wine_df, x="quality", y="total sulfur dioxide", jitter=True, ax=ax1)
sns.stripplot(data=white_wine_df, x="quality", y="free sulfur dioxide", jitter=True, ax=ax2)
```

Out[24]: <matplotlib.axes.\_subplots.AxesSubplot at 0x114f6d4e0>



```
In [25]: high_qual_ww_tsd_mean = white_wine_df[white_wine_df["quality"] >= 7]["total sulfur dioxide"].mean()
high_qual_ww_tsd_mean = format(high_qual_ww_tsd_mean, '.1f')
print(f"The mean for higher quality white wines (quality >= 7) is {high_qual_ww_tsd_mean}")
```

The mean for higher quality white wines (quality >= 7) is 125.2

## Interpretation White Wines


Both plots show, that higher quality white wines tend to have less total sulfur dioxide in it.

## red wine



```
In [26]: fig, (ax1, ax2) = plt.subplots(1,2)
fig.set_size_inches(14.5, 4.5)
fig.dpi = 300
sns.stripplot(data=red_wine_df, x="quality", y="total sulfur dioxide", jitter=True, ax=ax1)
sns.stripplot(data=red_wine_df, x="quality", y="free sulfur dioxide", jitter=True, ax=ax2)
```

Out[26]: <matplotlib.axes.\_subplots.AxesSubplot at 0x114e00cc0>



```
In [27]: high_qual_rw_tsd_mean = red_wine_df[red_wine_df["quality"] >= 7]["total sulfur dioxide"].mean()
high_qual_rw_tsd_mean = format(high_qual_rw_tsd_mean, '.1f')
print(f"The mean for higher quality red wines (quality >= 7) is {high_qual_rw_tsd_mean}")
```

The mean for higher quality red wines (quality >= 7) is 34.9

## Interpretation Red Wines

For the red wines, there are much lower concentrations of sulfur dioxides. Additionally, there seems to be no direct correlation between sulfur dioxide concentration and perceived quality.


## Conclusion: Sulfur Dioxides and Quality

Regarding high quality white wines ( $\geq 7$ ), those wines have a mean of sulfur dioxides of around 125. Respectively high quality Red Wines ( $\geq 7$ ) have a mean concentration of sulfur dioxide of 35.

## Sulphates and Quality

```
In [28]: sns.stripplot(data=wine_df, x="quality", y="sulphates", jitter=True, hue="color", split=True)
```

Out[28]: <matplotlib.axes.\_subplots.AxesSubplot at 0x117849240>



## Alcohol in Wine

### White Wine

```
In [29]: white_wine_df.groupby("quality")["alcohol"].describe()
```

Out[29]:

|         | count  | mean      | std      | min  | 25%   | 50%   | 75%   | max  |
|---------|--------|-----------|----------|------|-------|-------|-------|------|
| quality |        |           |          |      |       |       |       |      |
| 3       | 20.0   | 10.345000 | 1.224089 | 8.0  | 9.55  | 10.45 | 11.00 | 12.6 |
| 4       | 163.0  | 10.152454 | 1.003217 | 8.4  | 9.40  | 10.10 | 10.75 | 13.5 |
| 5       | 1457.0 | 9.808840  | 0.847065 | 8.0  | 9.20  | 9.50  | 10.30 | 13.6 |
| 6       | 2198.0 | 10.575372 | 1.147776 | 8.5  | 9.60  | 10.50 | 11.40 | 14.0 |
| 7       | 880.0  | 11.367936 | 1.246536 | 8.6  | 10.60 | 11.40 | 12.30 | 14.2 |
| 8       | 175.0  | 11.636000 | 1.280138 | 8.5  | 11.00 | 12.00 | 12.60 | 14.0 |
| 9       | 5.0    | 12.180000 | 1.013410 | 10.4 | 12.40 | 12.50 | 12.70 | 12.9 |

### Red Wine

```
In [30]: red_wine_df.groupby("quality")["alcohol"].describe()
```


```
Out[30]:
```

|         | count | mean      | std      | min | 25%    | 50%    | 75%    | max  |
|---------|-------|-----------|----------|-----|--------|--------|--------|------|
| quality |       |           |          |     |        |        |        |      |
| 3       | 10.0  | 9.955000  | 0.818009 | 8.4 | 9.725  | 9.925  | 10.575 | 11.0 |
| 4       | 53.0  | 10.265094 | 0.934776 | 9.0 | 9.600  | 10.000 | 11.000 | 13.1 |
| 5       | 681.0 | 9.899706  | 0.736521 | 8.5 | 9.400  | 9.700  | 10.200 | 14.9 |
| 6       | 638.0 | 10.629519 | 1.049639 | 8.4 | 9.800  | 10.500 | 11.300 | 14.0 |
| 7       | 199.0 | 11.465913 | 0.961933 | 9.2 | 10.800 | 11.500 | 12.100 | 14.0 |
| 8       | 18.0  | 12.094444 | 1.224011 | 9.8 | 11.325 | 12.150 | 12.875 | 14.0 |

## Plotting Alcohol to Quality


```
In [31]: sns.barplot(data=wine_df, x="quality", y="alcohol", hue="color")
```

```
Out[31]: <matplotlib.axes._subplots.AxesSubplot at 0x114e342e8>
```



```
In [32]: sns.lmplot(data=wine_df, x="quality", y="alcohol", hue="color")
```

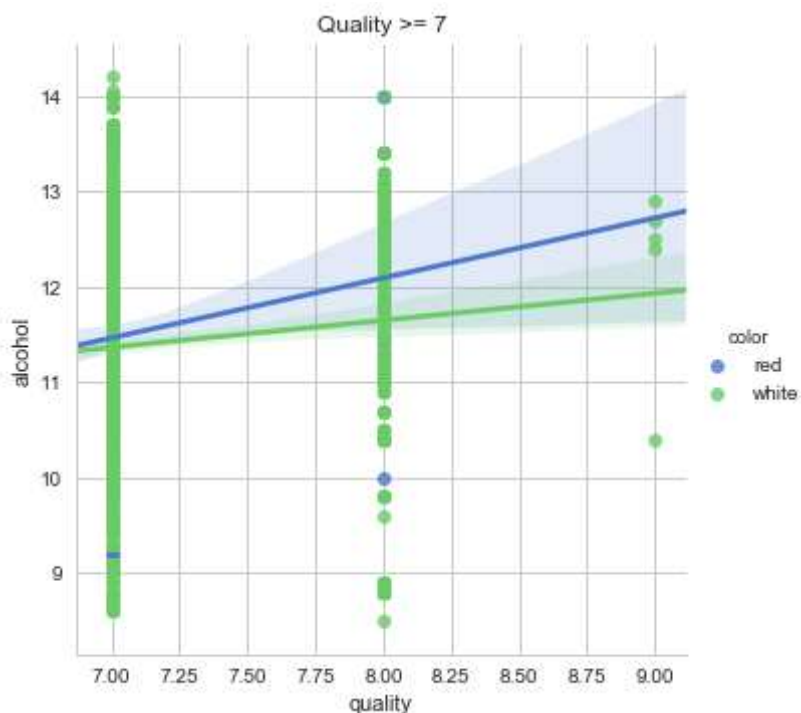
```
Out[32]: <seaborn.axisgrid.FacetGrid at 0x118bd3160>
```



### Alcohol to Quality relation for Wines equal or greater than 7

```
In [33]: hq_wines = wine_df[wine_df.quality >= 7]
sns.lmplot(data=hq_wines, x="quality", y="alcohol", hue="color")
sns.plt.title("Quality >= 7")
```


```
Out[33]: <matplotlib.text.Text at 0x118fe82e8>
```



### Heatmap Alcohol to Quality

```
In [34]: heat_table = wine_df[["quality", "alcohol"]].copy()
heat_table["alcohol"] = heat_table.alcohol.apply(func=lambda x: round(x * 2) / 2)
heat_table = heat_table.groupby(["quality", "alcohol"])["alcohol"].count().reset_index(name='counts')
sns.heatmap(heat_table.pivot("quality", "alcohol", "counts"))
```


Out[34]: <matplotlib.axes.\_subplots.AxesSubplot at 0x118e0eb70>



## Chlorides

```
In [35]: sns.barplot(data=wine_df, hue="color", x="quality", y="chlorides")
```

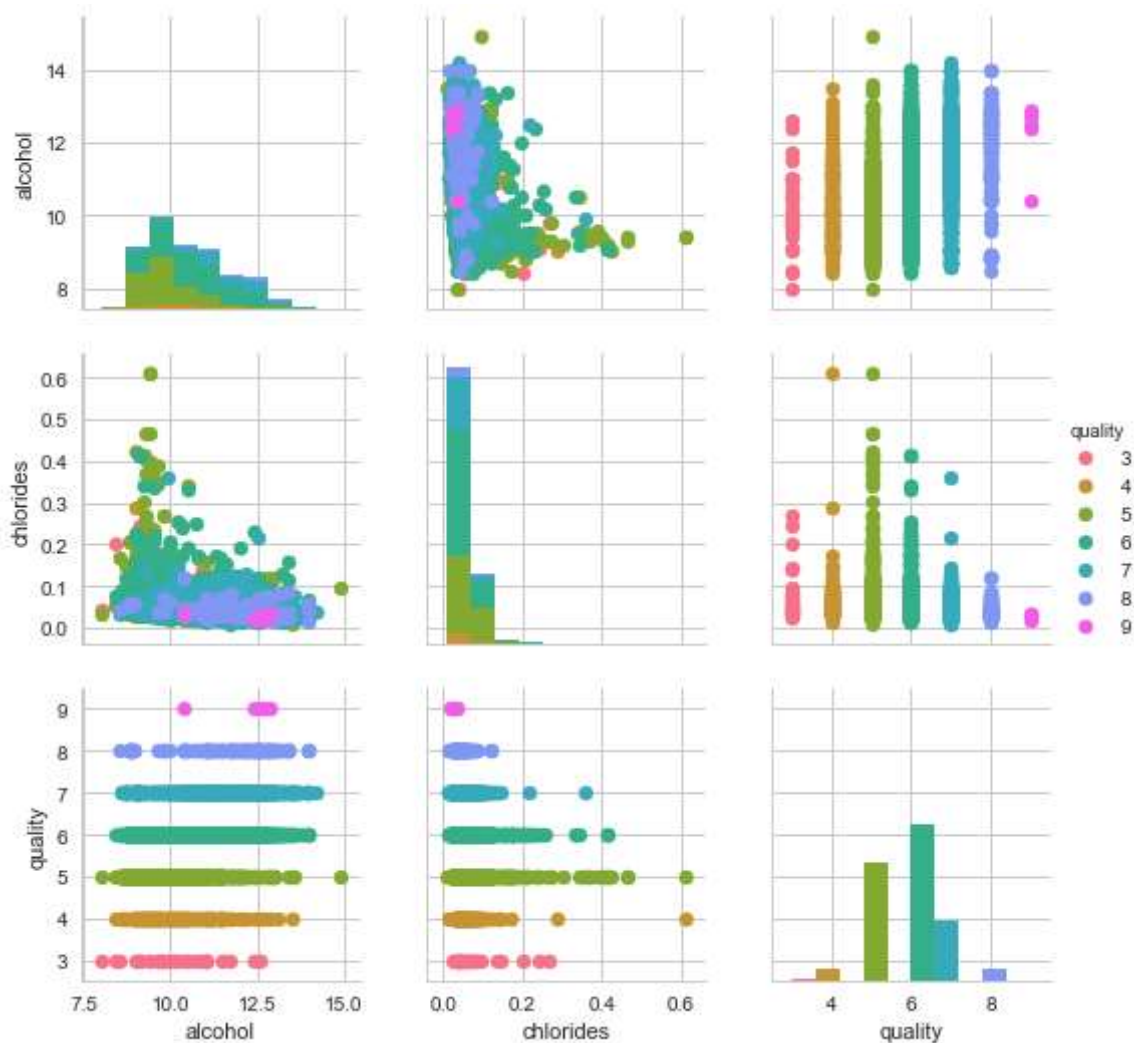
Out[35]: <matplotlib.axes.\_subplots.AxesSubplot at 0x1192a8c18>



**The less chlorides in a wine the higher the quality.**

# Chlorides and Alcohol

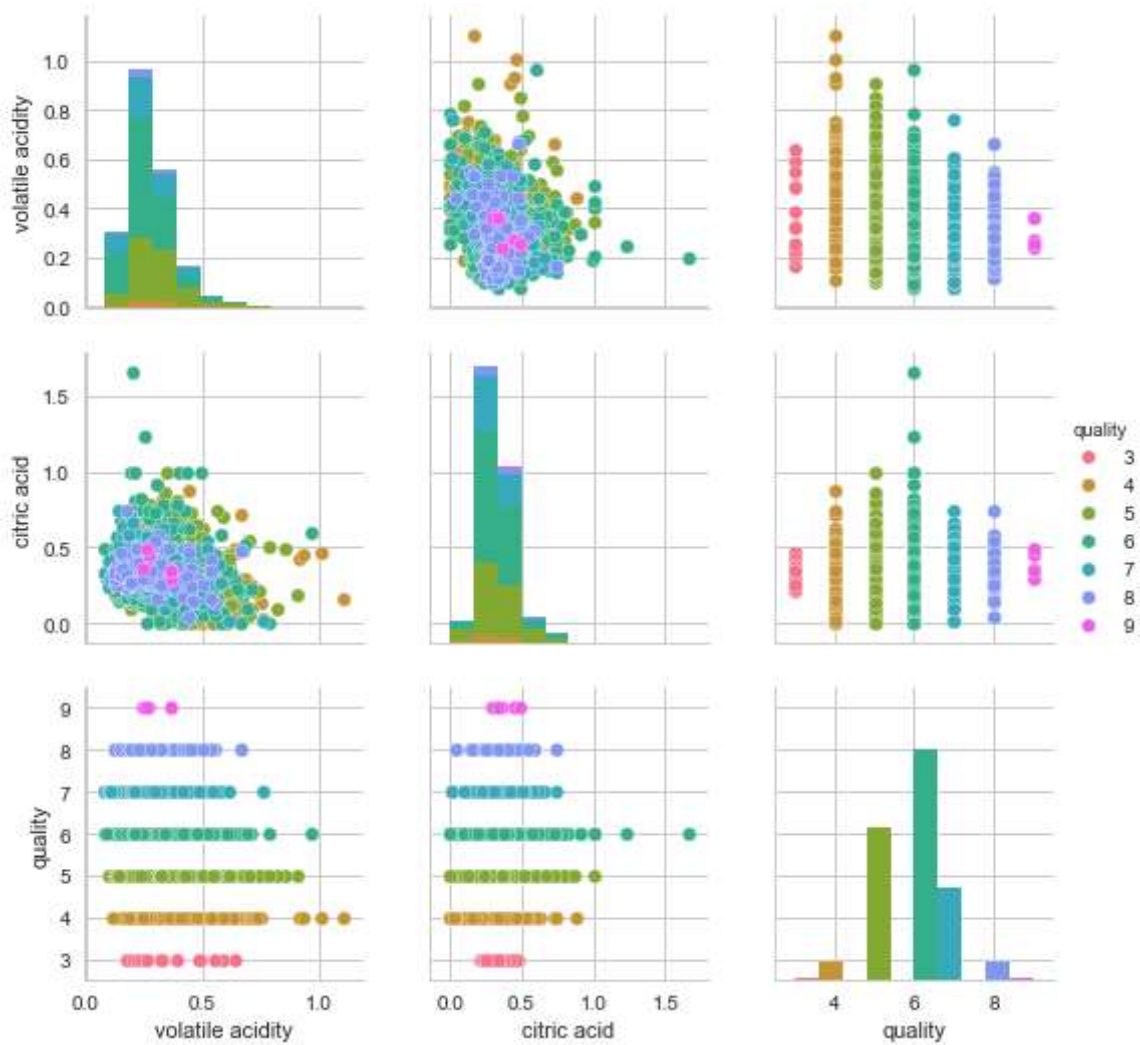
```
In [36]: g = sns.PairGrid(wine_df[["alcohol", "chlorides", "quality"]], hue="quality")
g = g.map_diag(plt.hist)
g = g.map_offdiag(plt.scatter)
g = g.add_legend()
```



## Acids

```
In [37]: sns.pairplot(white_wine_df[["volatile acidity", "citric acid", "quality"]], hue="quality")
```

```
Out[37]: <seaborn.axisgrid.PairGrid at 0x1194682b0>
```

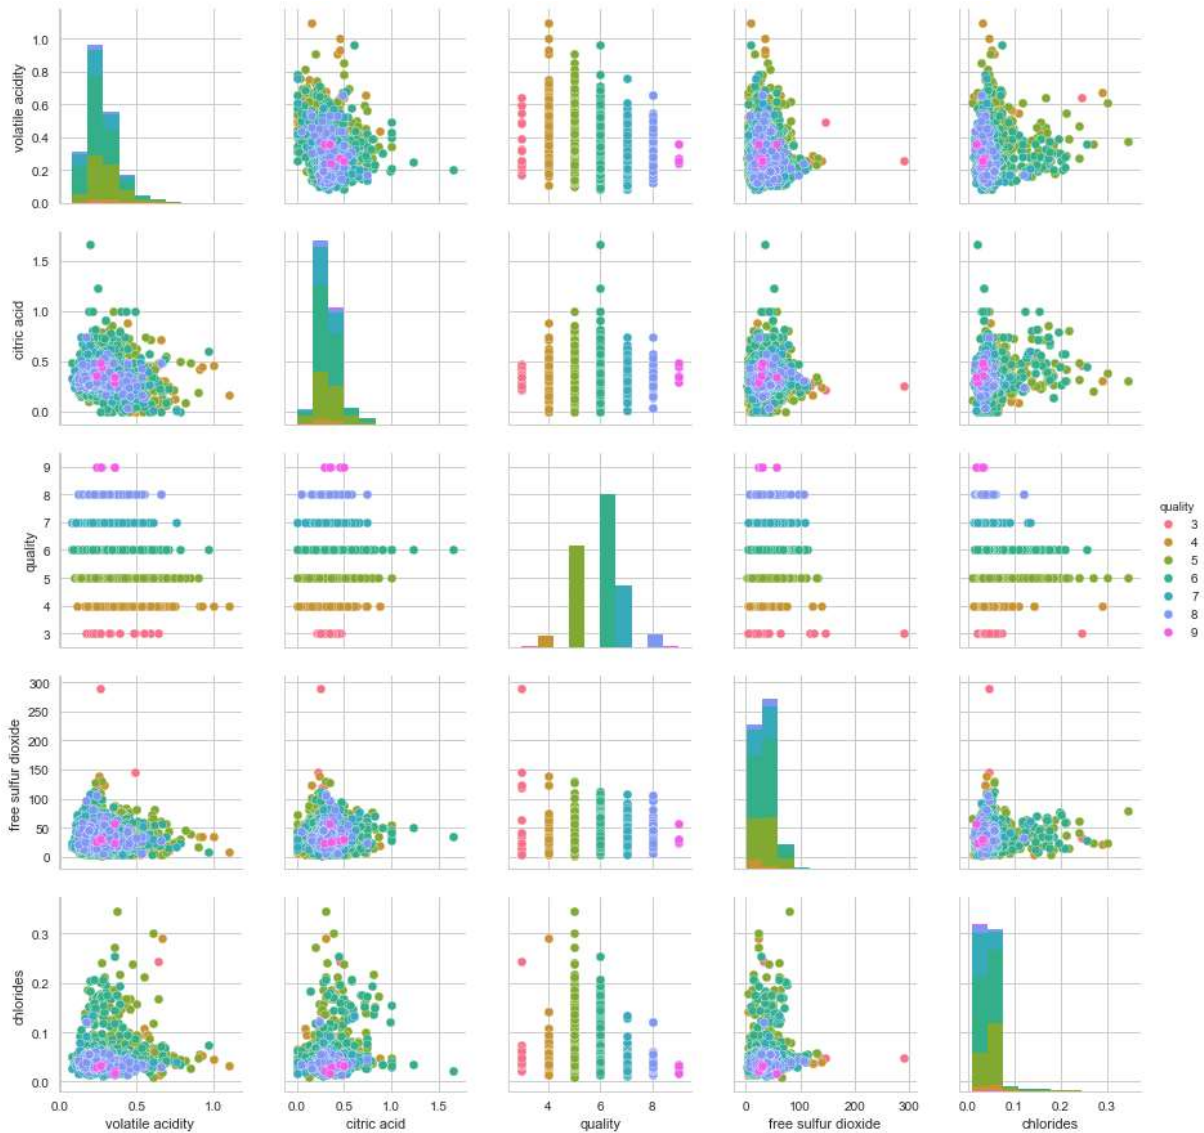


**Bringing the relevant attributes together**



```
In [38]: sns.pairplot(white_wine_df[["volatile acidity", "citric acid", "quality", "free sulfur dioxide", "chlorides"]], hue="quality")
```

```
Out[38]: <seaborn.axisgrid.PairGrid at 0x11a544128>
```



```
In [39]: white_wine_df[["volatile acidity", "citric acid", "quality"]].groupby("quality").describe(percentiles=[])
```

```
Out[39]:
```

|         | citric acid |          |          |      |       |      | volatile acidity |          |          |      |     |
|---------|-------------|----------|----------|------|-------|------|------------------|----------|----------|------|-----|
|         | count       | mean     | std      | min  | 50%   | max  | count            | mean     | std      | min  | 50% |
| quality |             |          |          |      |       |      |                  |          |          |      |     |
| 3       | 20.0        | 0.336000 | 0.081460 | 0.21 | 0.345 | 0.47 | 20.0             | 0.333250 | 0.140827 | 0.17 | 0.2 |
| 4       | 163.0       | 0.304233 | 0.163857 | 0.00 | 0.290 | 0.88 | 163.0            | 0.381227 | 0.173463 | 0.11 | 0.3 |
| 5       | 1457.0      | 0.337653 | 0.140814 | 0.00 | 0.320 | 1.00 | 1457.0           | 0.302011 | 0.100066 | 0.10 | 0.2 |
| 6       | 2198.0      | 0.338025 | 0.119325 | 0.00 | 0.320 | 1.66 | 2198.0           | 0.260564 | 0.088142 | 0.08 | 0.2 |
| 7       | 880.0       | 0.325625 | 0.079183 | 0.01 | 0.310 | 0.74 | 880.0            | 0.262767 | 0.091106 | 0.08 | 0.2 |
| 8       | 175.0       | 0.326514 | 0.085439 | 0.04 | 0.320 | 0.74 | 175.0            | 0.277400 | 0.108029 | 0.12 | 0.2 |
| 9       | 5.0         | 0.386000 | 0.082037 | 0.29 | 0.360 | 0.49 | 5.0              | 0.298000 | 0.057619 | 0.24 | 0.2 |



```
In [40]: white_wine_df[["quality", "free sulfur dioxide", "chlorides"]].groupby("quality").describe(percentiles=[])
```

Out[40]:

|         | chlorides |          |          |       |       |       | free sulfur dioxide |           |           |      |
|---------|-----------|----------|----------|-------|-------|-------|---------------------|-----------|-----------|------|
|         | count     | mean     | std      | min   | 50%   | max   | count               | mean      | std       | min  |
| quality |           |          |          |       |       |       |                     |           |           |      |
| 3       | 20.0      | 0.054300 | 0.046468 | 0.022 | 0.041 | 0.244 | 20.0                | 53.325000 | 69.420776 | 5.0  |
| 4       | 163.0     | 0.050098 | 0.025888 | 0.013 | 0.046 | 0.290 | 163.0               | 23.358896 | 20.391349 | 3.0  |
| 5       | 1457.0    | 0.051546 | 0.026496 | 0.009 | 0.047 | 0.346 | 1457.0              | 36.432052 | 18.145991 | 2.0  |
| 6       | 2198.0    | 0.045217 | 0.020453 | 0.015 | 0.043 | 0.255 | 2198.0              | 35.650591 | 15.735679 | 3.0  |
| 7       | 880.0     | 0.038191 | 0.010697 | 0.012 | 0.037 | 0.135 | 880.0               | 34.125568 | 13.244737 | 5.0  |
| 8       | 175.0     | 0.038314 | 0.013164 | 0.014 | 0.036 | 0.121 | 175.0               | 36.720000 | 16.203675 | 6.0  |
| 9       | 5.0       | 0.027400 | 0.007436 | 0.018 | 0.031 | 0.035 | 5.0                 | 33.400000 | 13.427584 | 24.0 |

## Final Conclusion

no attribute alone is strong enough to define a high quality wine, but as the figures show. For a wine to score high, having the acids and sulfur dioxide values all within in a certain range can help.

## Best vs. Worst

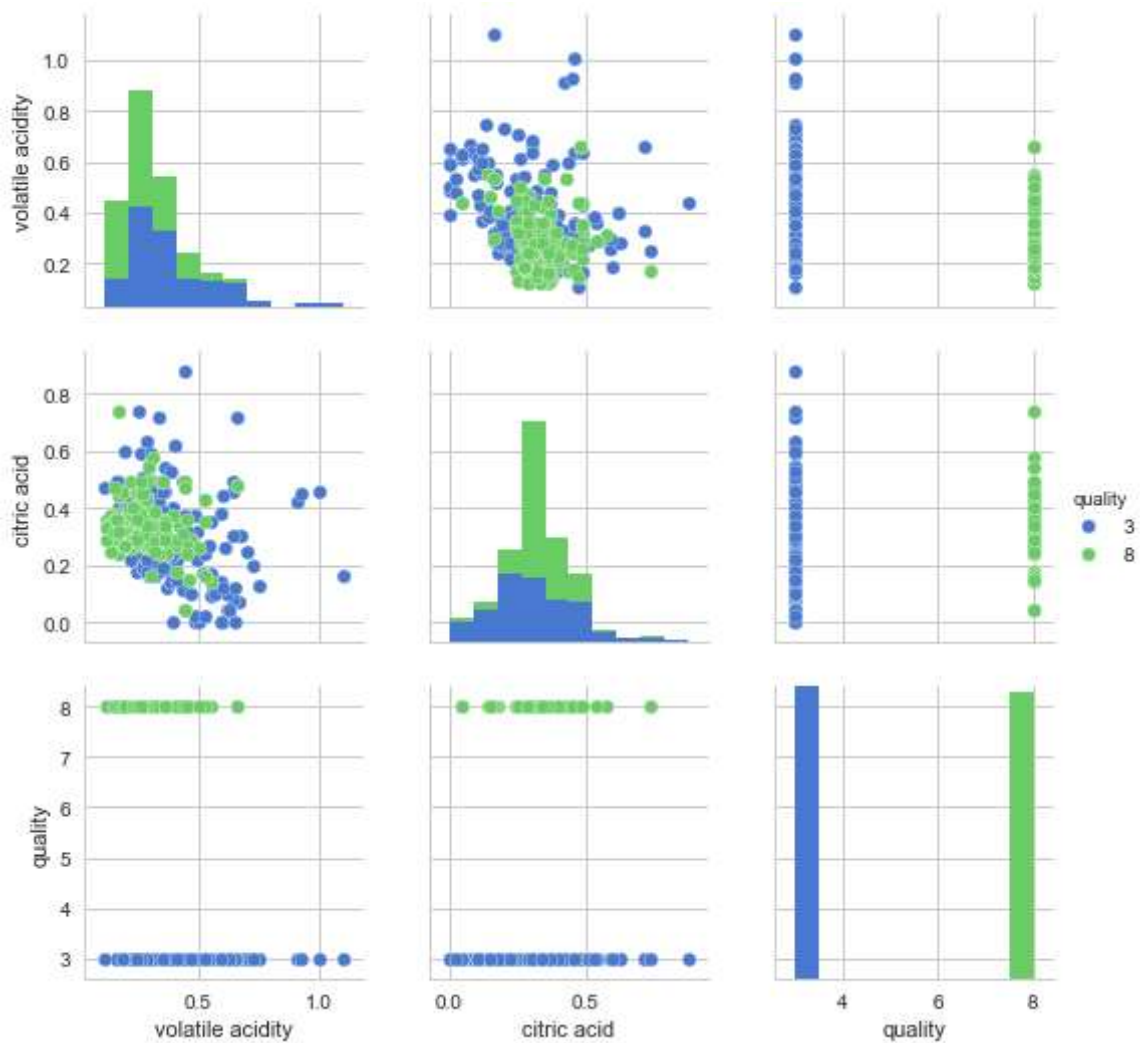
when comparing the best (8,9) vs. worst (3,4) we can see that they well overlap each other.

```

In [41]: qual3 = white_wine_df[white_wine_df["quality"] == 3]
qual4 = white_wine_df[white_wine_df["quality"] == 4].copy()
qual4.quality = 3
qual8 = white_wine_df[white_wine_df["quality"] == 8]
qual9 = white_wine_df[white_wine_df["quality"] == 9].copy()
qual9.quality = 8
white_wine_sample = pd.concat([qual3, qual4, qual8, qual9], ignore_index=True)
white_wine_sample
sns.pairplot(white_wine_sample[["volatile acidity", "citric acid", "quality"]], hue="quality")

```

Out[41]: <seaborn.axisgrid.PairGrid at 0x11cc89ef0>



```
In [42]: sns.pairplot(white_wine_sample[["volatile acidity", "citric acid", "quality", "sulphates", "chlorides"]], hue="quality")
```

```
Out[42]: <seaborn.axisgrid.PairGrid at 0x11d07cb00>
```



**Seems like labs can't measure a wine's inner spirit (yet).**

**But if you have to pick a wine only based on specs, i would suggest white wines close to this values:**

```
In [43]: qual8[["quality", "chlorides", "alcohol", "citric acid", "sulphates"]].describe(percentiles=[])
```

```
Out[43]:
```

|       | quality | chlorides  | alcohol    | citric acid | sulphates  |
|-------|---------|------------|------------|-------------|------------|
| count | 175.0   | 175.000000 | 175.000000 | 175.000000  | 175.000000 |
| mean  | 8.0     | 0.038314   | 11.636000  | 0.326514    | 0.486229   |
| std   | 0.0     | 0.013164   | 1.280138   | 0.085439    | 0.147073   |
| min   | 8.0     | 0.014000   | 8.500000   | 0.040000    | 0.250000   |
| 50%   | 8.0     | 0.036000   | 12.000000  | 0.320000    | 0.460000   |
| max   | 8.0     | 0.121000   | 14.000000  | 0.740000    | 0.950000   |

and red wines close to this values:

```
In [45]: rqual8 = red_wine_df[red_wine_df["quality"] >= 8]
rqual8[["quality", "chlorides", "alcohol", "citric acid", "sulphates"]].describe(percentiles=[])
```

Out[45]:

|       | quality | chlorides | alcohol   | citric acid | sulphates |
|-------|---------|-----------|-----------|-------------|-----------|
| count | 18.0    | 18.000000 | 18.000000 | 18.000000   | 18.000000 |
| mean  | 8.0     | 0.068444  | 12.094444 | 0.391111    | 0.767778  |
| std   | 0.0     | 0.011678  | 1.224011  | 0.199526    | 0.115379  |
| min   | 8.0     | 0.044000  | 9.800000  | 0.030000    | 0.630000  |
| 50%   | 8.0     | 0.070500  | 12.150000 | 0.420000    | 0.740000  |
| max   | 8.0     | 0.086000  | 14.000000 | 0.720000    | 1.100000  |



```
In [ ]:
```